


Minimization of Maximum Mean Discrepancy (MMD)

Zonghao (Hudson) Chen

Annual student conference, UCL Computer Science, 2026



Zonghao (Hudson) Chen

4th year PhD student, University College London

Foundational AI Centre
Gatsby Unit


François-Xavier Briol
Arthur Gretton

Tsinghua Uni., B.Eng. in Electronic Engineering, 2022

Research Interests


- Kernel and nonparametric methods
- Causal inference
- Statistical learning theory

A research timeline from the start of my PhD



Background: what is MMD?

- Given two distributions $\mathbb{P}$ and $\mathbb{Q}$, how to compare them?
- Compute the difference of their means: $|\mathbb{E}_{\mathbb{P}}[X] - \mathbb{E}_{\mathbb{Q}}[X]|$.
- Compute the difference of their second moments: $|\mathbb{E}_{\mathbb{P}}[X^2] - \mathbb{E}_{\mathbb{Q}}[X^2]|$.
- The first two moments are not enough!




Goal: Compare $\mathbb{P}$ and $\mathbb{Q}$ with a general f , i.e., $|\mathbb{E}_{\mathbb{P}}[f(X)] - \mathbb{E}_{\mathbb{Q}}[f(X)]|$?

Background: Maximum Mean Discrepancy (MMD) [GBR⁺12]

- $\mathcal{H}$ is a reproducing kernel Hilbert space (RKHS) with kernel k .

$$\text{MMD}(\mathbb{P}, \mathbb{Q}) := \sup_{f \in \mathcal{H}} |\mathbb{E}_{\mathbb{P}}[f(X)] - \mathbb{E}_{\mathbb{Q}}[f(X)]|.$$

- The RKHS contains functions that are linear combinations of the feature maps $\phi : \mathcal{X} \rightarrow \mathcal{H}$, $f = \sum_{i=1}^n \alpha_i \phi(x_i)$.



Theorem 1

MMD admits a closed-form expression:

$$\text{MMD}^2(\mathbb{P}, \mathbb{Q}) = \mathbb{E}_{X, X' \sim \mathbb{P}}[k(X, X')] + \mathbb{E}_{Y, Y' \sim \mathbb{Q}}[k(Y, Y')] - 2\mathbb{E}_{X \sim \mathbb{P}, Y \sim \mathbb{Q}}[k(X, Y)].$$




- MMD is a probability divergence ($\text{MMD}(\mathbb{P}, \mathbb{Q}) = 0 \iff \mathbb{P} = \mathbb{Q}$) if k is characteristic (Gaussian) [SFL11].

Background: Maximum Mean Discrepancy (MMD)

We now have MMD

A kernel-based way to compare probability distributions $\mathbb{P}$ and $\mathbb{Q}$.

What can we use it for?




Numerical Integration

Part I: Numerical Integration


Where do integrals appear in computer science?



Computer graphics
Rendering



Bayesian inference
Posterior expectations



Reinforcement learning
Expected returns

Part I: Numerical Integration

- Given a function f and a distribution $\mathbb{P}$, how to compute $\int f(x)d\mathbb{P}(x)$?
- Given a set of points $\{x_i\}_{i=1}^n$ and function evaluations $\{f(x_i)\}_{i=1}^n$.
- Quadrature rule:

$$\hat{I} = \sum_{i=1}^n w_i f(x_i).$$

- When $w_1 = \dots = w_n = 1/n$, this is Monte Carlo. $\hat{I}_{MC} = \frac{1}{n} \sum_{i=1}^n f(x_i)$.

$$\mathbb{E} \left[|\hat{I}_{MC} - I| \right] = \mathcal{O}(n^{-1/2}).$$

- Or equivalently, to obtain an ϵ -accurate estimate, we need $n = \mathcal{O}(\epsilon^{-2})$ samples.

Question: Can we do better than $\mathcal{O}(n^{-1/2})$?

Answer: Find better weights $w_1, \dots, w_n$.

Part I: Numerical Integration

- Suppose $f \in \mathcal{H}$ is in the RKHS associated with k .
- Denote $\hat{\mathbb{P}} = \sum_{i=1}^n w_i \delta_{x_i}$ as the weighted empirical distribution.

$$\left| I - \hat{I} \right|^2 = \left| \int f(x) d\mathbb{P}(x) - \sum_{i=1}^n w_i f(x_i) \right|^2 \leq \|f\|_{\mathcal{H}}^2 \cdot \text{MMD}^2(\hat{\mathbb{P}}, \mathbb{P}).$$

- The right-hand side is the **worst-case** error of the quadrature rule.
- The **optimal** weights that minimize $\text{MMD}(\hat{\mathbb{P}}, \mathbb{P})$:


$$w_{1:n} = \mathbb{E}_{X \sim \mathbb{P}}[k(X, x_{1:n})] K(x_{1:n}, x_{1:n})^{-1}.$$

- The kernel quadrature (KQ) rule: $\hat{I}_{\text{KQ}} = \sum_{i=1}^n w_i f(x_i)$.


Part I: Numerical Integration

Rate: KQ achieves $\mathcal{O}(n^{-s/d})$ rate.

- s is the smoothness of f , d is the dimension of input x .
- Better than $\mathcal{O}(n^{-1/2})$ when $s > d/2$.




Part I: Numerical Integration



- $s = 4$ and $d = 1$.

Part I: Numerical Integration

- Nested expectation: $I := \mathbb{E}_{\theta \sim \mathbb{Q}} [f(\mathbb{E}_{X \sim \mathbb{P}_\theta} [g(X, \theta)])]$



Total simulator evaluations: $N \times T$

Stage I: inner KQ

$$\hat{J}_{\text{KQ}}(\theta_t) = \sum_{n=1}^N w_{n,t}^{\mathcal{X}} g(x_n^{(t)}, \theta_t)$$

Stage II: outer KQ

$$\hat{I}_{\text{NKQ}} = \sum_{t=1}^T w_t^{\Theta} f(\hat{J}^{\text{KQ}}(\theta_t))$$

Cost to achieve error $|I - \hat{I}| \leq \Delta$:

NMC: $\mathcal{O}(\Delta^{-4})$ or $\mathcal{O}(\Delta^{-3})$

MLMC: $\mathcal{O}(\Delta^{-2})$

NQMC: $\mathcal{O}(\Delta^{-2.5})$

NKQ: $\tilde{\mathcal{O}}(\Delta^{-d_{\mathcal{X}}/s_{\mathcal{X}} - d_{\Theta}/s_{\Theta}})$

Conditional Bayesian Quadrature (UAI 2025)

Z. Chen, M. Naslidnyk, A. Gretton, F.-X. Briol

- Extension to conditional expectations: $I(\theta) = \mathbb{E}_{X \sim p(\cdot|\theta)}[g(X, \theta)]$.
- Weighted rule: $\hat{I}_{\text{CBQ}}(\theta) := \sum_{t=1}^T \sum_{n=1}^N w_{n,t} g(x_n^{(t)}, \theta_t)$.

Nested Expectations with Kernel Quadrature (ICML 2025)

Z. Chen, M. Naslidnyk, F.-X. Briol

- Extension to nested expectations: $I := \mathbb{E}_{\theta \sim \mathbb{Q}} [f(\mathbb{E}_{X \sim \mathbb{P}_\theta} [g(X, \theta)])]$.
- Weighted rule: $\hat{I}_{\text{NKQ}} = \sum_{t=1}^T w_t^\Theta f\left(\sum_{n=1}^N w_{n,t}^\chi g(x_n^{(t)}, \theta_t)\right)$.

BayesSum: Bayesian Quadrature in Discrete Spaces (ProbNum 2026)


S. S. Kang, F.-X. Briol, T. Karvonen, Z. Chen

- Extension to sums: $I = \mathbb{E}_{X \sim \mathbb{P}}[f(X)] = \sum_{x \in \mathcal{X}} f(x)p(x)$.

Sampling / Generative Modelling


Quadrature

fixed locations, choose weights




Sampling / generation

equal weights, choose locations




Part II: Sampling / Generative modelling

Sampling



Generative modelling



- Given a target distribution $\mathbb{Q}$, how to draw samples $\{x_i\}_{i=1}^n$ from it?
 - When $\mathbb{Q}$ is known up to normalization, this is sampling.
 - When $\mathbb{Q}$ is known with i.i.d samples, this is generative modelling.

Problem: How to learn a target probability distribution $\mathbb{Q}$ on $\mathbb{R}^d$.

- This problem can be written as an **optimization** problem on $\mathcal{P}(\mathbb{R}^d)$, the space of probability distributions.

$$\arg \min_{\mathbb{P} \in \mathcal{P}(\mathbb{R}^d)} D(\mathbb{P}, \mathbb{Q}).$$


- Here D is a similarity metric or distance, e.g. Kullback–Leibler divergence / MMD.
 - $D(\mathbb{P}_1, \mathbb{P}_2) = 0$ if and only if $\mathbb{P}_1 = \mathbb{P}_2$.
- How to find the minimum? Gradient descent!

Part II: Sampling / Generative modelling


Question: Wait... How to define gradients on $\mathcal{P}(\mathbb{R}^d)$?

Answer: Wasserstein gradient ! [Ott01]

Optimal transport



Wasserstein gradient




$$\nabla_W \mathcal{F}(\mathbb{P}_0) = \lim_{\varepsilon \rightarrow 0} \frac{\mathcal{F}(\mathbb{P}_\varepsilon) - \mathcal{F}(\mathbb{P}_0)}{\varepsilon}$$

Part II: Sampling / Generative modelling

- Unsurprisingly, we use MMD as the divergence.

$$\arg \min_{\mathbb{P} \in \mathcal{P}(\mathbb{R}^d)} \text{MMD}^2(\mathbb{P}, \mathbb{Q}).$$

- A (Wasserstein) gradient descent of $\text{MMD}^2(\mathbb{P}, \mathbb{Q})$ gives a sequence of distributions: $\mathbb{P}_0, \mathbb{P}_1, \dots, \mathbb{P}_t$ [AKSG19].
 - In the limit, $\mathbb{P}_t$ converges to $\mathbb{Q}$ and $\text{MMD}(\mathbb{P}_t, \mathbb{Q}) \rightarrow 0$ (or not?).



- MMD flow generates more **uniform** samples.

Part II: Sampling / Generative modelling



- MMD flow does not outperform diffusion-based generative models [GBG25].
- A novel framework of generative modelling & sampling.

Challenge: Non-convexity of the squared MMD objective: $\lim_{t \rightarrow \infty} \text{MMD}^2(\mathbb{P}_t, \mathbb{Q}) > 0$.

- $\chi^2(\mu \parallel \pi)$ is convex in μ [OT11], yet hard to implement.
- DrMMD *interpolates* between MMD ($\lambda = \infty$) and χ^2 ($\lambda = 0$).

$$\text{DrMMD}(\mu, \pi) = \sup_{f \in \mathcal{H}: \|f\|_{L_2(\mu)}^2 + \lambda \|f\|_{\mathcal{H}}^2 \leq 1} \int f \, d(\mu - \pi).$$

- DrMMD flow converges in KL divergence: $\lim_{t \rightarrow \infty} \text{KL}(\mathbb{P}_t, \mathbb{Q}) = 0$.
 - When π is ‘Gaussian-like’.

(De)-regularized MMD Gradient Flow (JMLR 2025)

Z. Chen, A. Mustafi, P. Glaser, A. Korba, A. Gretton, B. K. Sriperumbudur

- DrMMD interpolates between MMD and χ^2 divergence.
- Convergence in KL divergence: $\lim_{t \rightarrow \infty} \text{KL}(\mathbb{P}_t, \mathbb{Q}) = 0$.

Sobolev Regularized MMD Gradient Flow (Arxiv)

C. Tian, B. K. Sriperumbudur, A. Gretton, Z. Chen


- Gradient regularization on the critic [MSR19].

$$\text{SrMMD}(\mu, \pi) = \sup_{f \in \mathcal{H}: \|\nabla f\|_{L^d_2(\mu)}^2 + \lambda \|f\|_{\mathcal{H}}^2 \leq 1} \int f \, d(\mu - \pi).$$

- Convergence in MMD: $\lim_{t \rightarrow \infty} \text{MMD}(\mathbb{P}_t, \mathbb{Q}) = 0$.
- No conditions on π .

Empirical observation:

(Wasserstein) gradient flows give samples more **uniform** than i.i.d samples [XKS22].



- However, theoretical rate can hardly beat Monte Carlo: $n^{-1/2}$.

Part II: Sampling / Generative modelling

- For any $i \in \{1, \dots, n\}$:

$$\nabla_{x_i} \text{MMD}^2 \left(\mu, \frac{1}{n} \sum_{j=1}^n \delta_{x_j} \right) = 0 \iff \frac{1}{n} \sum_{j=1}^n \nabla_1 k(x_i, x_j) - \int \nabla_1 k(x_i, y) d\mu(y) = 0.$$

- Exact integration on $x \mapsto \nabla_1 k(x_i, x)$.

Exact integration subspace:

$$\mathcal{G}_n := \text{span} \left\{ \nabla_1 k(x_i, \cdot) : i \in \{1, \dots, n\} \right\} \subset \mathcal{H}.$$

- As $n \rightarrow \infty$, $\mathcal{G}_n$ approximates the whole RKHS $\mathcal{H}$ up to constant functions.
- Stationary MMD points achieve $\mathcal{O}(n^{-1/2})$ rate for numerical integration.

Stationary MMD Points (ICML 2026)

Z. Chen, T. Karvonen, H. Kanagawa, F.-X. Briol, C. J. Oates

- Samples from MMD flow are better for numerical integration than IID samples.
- IID samples: $\mathcal{O}(n^{-1/2})$, stationary MMD points: $o(n^{-1/2})$.

Thinned Mean Field Langevin Dynamics (ICML 2026)

Z. Chen, H. Kanagawa, F.-X. Briol, C. J. Oates, L. Mackey

- Computational cost of MMD flow is $\mathcal{O}(n^2)$ due to pair-wise interaction.
- Reduce cost to $\mathcal{O}(n^{1.5})$ without sacrificing the sample quality.
- Kernel thinning!

Statistical Inference




MMD flow

$$\arg \min_{\mathbb{P} \in \mathcal{P}(\mathbb{R}^d)} \text{MMD}^2(\mathbb{P}, \mathbb{Q})$$

Minimum MMD inference

$$\theta_{\text{MMD}} = \arg \min_{\theta \in \Theta} \text{MMD}^2(\mathbb{P}_{\theta}, \mathbb{Q})$$

Where does statistical inference appear in computer science?



Part III: Statistical inference


- Given a parametric model $\{\mathbb{P}_\theta\}_{\theta \in \Theta \subseteq \mathbb{R}^p}$.
- Infer the parameter θ that best explains an unknown data distribution $\mathbb{Q}$.
- Maximum likelihood estimation is the most popular:

$$\theta_{\text{MLE}} = \arg \max_{\theta \in \Theta} \mathbb{E}_{X \sim \mathbb{Q}} [\log p_\theta(X)].$$

What is not captured by maximum likelihood?


Likelihood-free simulators

When $\mathbb{P}_\theta$ is a scientific simulator: we can sample $x \sim p_\theta(x)$, but cannot evaluate the density $p_\theta(x)$.



Robustness

A few contaminated observations can strongly move the fitted parameter.



- Minimum MMD estimator [? ?]

$$\begin{aligned}\theta_{\text{MMD}} &= \arg \min_{\theta \in \Theta} \text{MMD}^2(\mathbb{P}_\theta, \mathbb{Q}) \\ &= \arg \min_{\theta \in \Theta} \left\{ \mathbb{E}_{x, x' \sim \mathbb{P}_\theta} [k(x, x')] - 2\mathbb{E}_{x \sim \mathbb{P}_\theta, y \sim \mathbb{Q}} [k(x, y)] + \text{const} \right\}.\end{aligned}$$

Challenge 1: Likelihood-free inference?

✓ Only requires samples from $\mathbb{P}_\theta$ and $\mathbb{Q}$.

Challenge 2: Robustness to outliers?

✓ Bounded effects under bounded kernels.

Question: How to reliably find θ_{MMD} ?


- MMD flow:

$$\mathbb{P}_{t+1} = (\text{Id} - \gamma \nabla f_{\mathbb{P}_t, \mathbb{Q}})_{\#} \mathbb{P}_t.$$

- MMD flow converges globally.
 - Adaptive lengthscale, DrMMD, SrMMD, etc.
- Consider $\mathbb{P}_{\theta} = (T_{\theta})_{\#} \rho$:

$$\mathbb{P}_{\theta_{t+1}} = (T_{\theta_t - \gamma \Delta \theta_t})_{\#} \rho.$$

- Choose $\mathbb{P}_{\theta_{t+1}}$ closest to the MMD flow iterate $\mathbb{P}_{t+1}$.
- $\lim_{t \rightarrow \infty} \text{MMD}(\mathbb{P}_{\theta_t}, \mathbb{Q}) = \min_{\theta} \text{MMD}_{\ell}(\mathbb{P}_{\theta}, \mathbb{Q})$.



On Optimization of Minimum MMD Estimators ([Arxiv](#))

S. S. Kang, L. Sharrock, X. Cheng, F.-X. Briol, [Z. Chen](#)

- A preconditioned method / natural gradient type method.
- Adaptive kernel lengthscale.

What next: causal inference.

Part IV: What next?

Minimum MMD for counterfactual inference

Question: What would have happened under a different intervention?

- Treatment A .
- Counterfactual outcome: $Y(a)$ under treatment $A = a$.
- Missing data/inference problem.



Conformal Counterfactual Inference under Hidden Confounding (KDD 2024)

[Z. Chen](#), R. Guo, J.-F. Ton, Y. Liu

Neural Networks for NPIV: Optimization and Generalization (Arxiv)

[Z. Chen](#), A. Nitanda, A. Gretton, T. Suzuki

NPIV with Observed Covariates (Arxiv)

Z. Shen, [Z. Chen](#), D. Meunier, I. Steinwart, A. Gretton, Z. Li

Summary

- MMD is a powerful tool for comparing probability distributions.

$$\text{MMD}(\mathbb{P}, \mathbb{Q}) := \sup_{f \in \mathcal{H}} |\mathbb{E}_{\mathbb{P}}[f(X)] - \mathbb{E}_{\mathbb{Q}}[f(X)]|.$$

- Minimization of MMD in numerical integration, sampling/generative modelling, and statistical inference.
- Future work: causal inference?

- [AKSG19] Michael Arbel, Anna Korba, Adil Salim, and Arthur Gretton. Maximum mean discrepancy gradient flow. In H. Wallach, H. Larochelle, A. Beygelzimer, F. d'Alché Buc, E. Fox, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc., 2019.
- [GBG25] Alexandre Galashov, Valentin De Bortoli, and Arthur Gretton. Deep MMD gradient flow without adversarial training. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [GBR⁺12] Arthur Gretton, Karsten M. Borgwardt, Malte J. Rasch, Bernhard Schölkopf, and Alexander Smola. A kernel two-sample test. *Journal of Machine Learning Research*, 13(1):723–773, 2012.
- [MSR19] Youssef Mroueh, Tom Sercu, and Anant Raj. Sobolev descent. In *The 22nd International Conference on Artificial Intelligence and Statistics*, pages 2976–2985. PMLR, 2019.
- [OT11] Shin-ichi Ohta and Asuka Takatsu. Displacement convexity of generalized relative entropies. *Advances in Mathematics*, 228(3):1742–1787, 2011.
- [Ott01] Felix Otto. The geometry of dissipative evolution equations: the porous medium equation. 2001.

- [SFL11] Bharath K. Sriperumbudur, Kenji Fukumizu, and Gert R. G. Lanckriet. Universality, characteristic kernels and RKHS embedding of measures. *Journal of Machine Learning Research*, 12(7):2389–2410, 2011.
- [XKS22] Lantian Xu, Anna Korba, and Dejan Slepcev. Accurate quantization of measures via interacting particle-based optimization. In *International Conference on Machine Learning*, pages 24576–24595. PMLR, 2022.